

Deep Learning

MSDS 631

Introduction to this course and
Deep Learning

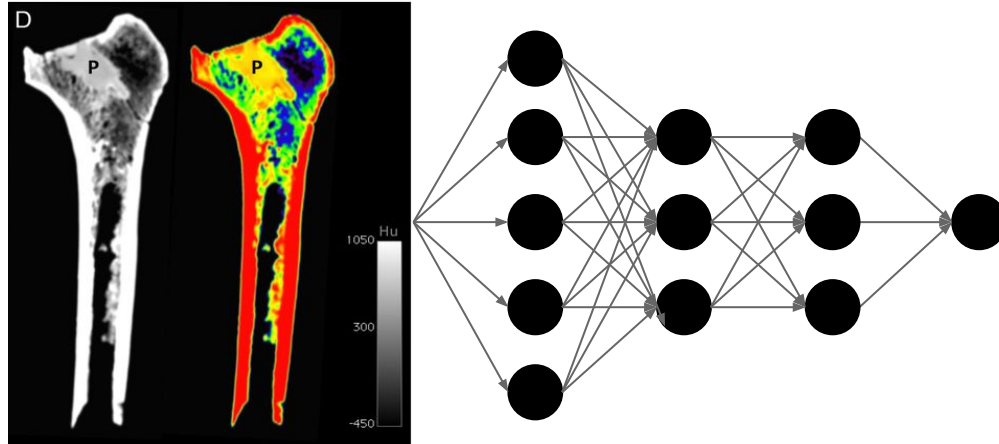
Michael Ruddy

Overview

- What we will cover and why
- How this course will operate
- What is Deep Learning really?
- How do I train a neural network in Pytorch?

What is Deep Learning?

- **Deep Learning (DL)** is a subset of Machine Learning where algorithms perform tasks by extracting *high-level features* from datasets that are usually very large and unstructured.
- Models are usually based on **artificial neural networks (ANNs or NNs)**.
 - Deep here refers to ANNs with many *layers*.

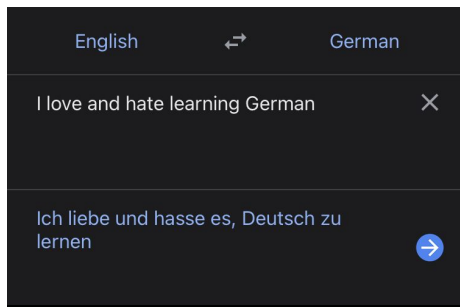


Why Deep Learning?

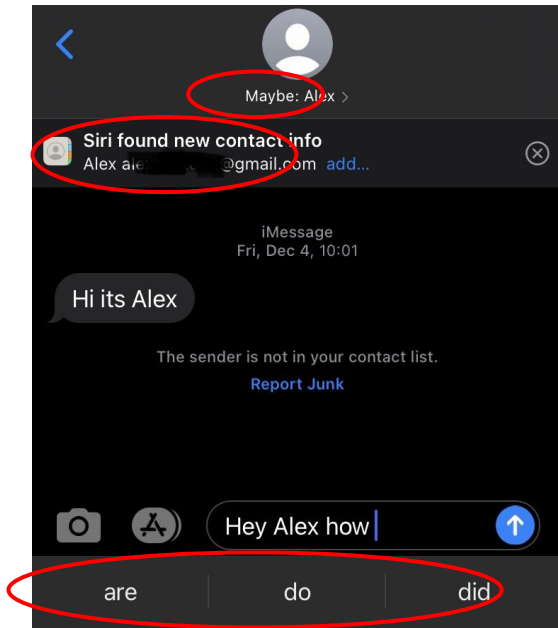
- Explosion in amount of data available and in computing power
 - Neural networks are often complicated models with many parameters, necessitating a lot of data and a lot of computing power
- Increasingly important aspect of data science
 - Image Science
 - Natural Language Processing

Why Deep Learning?

- Explosion in amount of data available and in computing power
 - Neural networks are often complicated models with many parameters, necessitating a lot of data and a lot of computing power
- Increasingly important aspect of data science
 - Image Science
 - Natural Language Processing



Some or all of the content shared in this Tweet conflicts with guidance from public health experts regarding COVID-19. [Learn more](#) [View](#)



What we will learn

- How to effectively train neural networks using PyTorch
 - Learning rates, batch norm, regularization, data augmentation, transfer learning, ...



What we will learn

- How to effectively train neural networks using PyTorch
 - Learning rates, batch norm, regularization, data augmentation, transfer learning, ...
- Image data
 - Common imaging tasks (classification, segmentation, ...)
 - Architectures for spatial data: CNNs
- Text data
 - Common NLP tasks (classification, comparison, ...)
 - Architectures for sequences/text (RNNs, Attention)

What we will learn

- How to effectively train neural networks using PyTorch
 - Learning rates, batch norm, regularization, data augmentation, transfer learning, ...
- Image data
 - Common imaging tasks (classification, segmentation, ...)
 - Architectures for spatial data: CNNs
- Text data
 - Common NLP tasks (classification, comparison, ...)
 - Architectures for sequences/text (RNNs, Attention)
- Maybe, if time/interest
 - GANs (style transfer, synthetic images/text)
 - Something else?

Disclaimer

- Too much in Deep Learning to go over everything!
- Goals
 - Provide a good foundation for whatever most interests you
 - Good mix of conceptual understanding and implementation in PyTorch

Disclaimer

- Too much in Deep Learning to go over everything!
- Goals
 - Provide a good foundation for whatever most interests you
 - Good mix of conceptual understanding and implementation in PyTorch
- Changes every year (sometimes multiple times)
 - What techniques work the best
 - Why a certain technique works
 - What is really going on in Deep Learning
 - What is the hot stuff everyone wants to do

Disclaimer

- Too much in Deep Learning to go over everything!
- Goals
 - Provide a good foundation for whatever most interests you
 - Good mix of conceptual understanding and implementation in PyTorch
- Changes every year (sometimes multiple times)
 - What techniques work the best
 - Why a certain technique works
 - What is really going on in Deep Learning
 - What is the hot stuff everyone wants to do
- Don't accept a research paper as the absolute truth

How this class will run: What you'll do

- On Tuesday/Thursday: **Live Lecture**
 - Quizzes graded for completion (Polls)
 - Active participation (answer if called on)
 - In-class work (group optional; determined before class)

How this class will run: What you'll do

- On Tuesday/Thursday: **Live Lecture**
 - Quizzes graded for completion (Polls)
 - Active participation (answer if called on)
 - In-class work (group optional; determined before class)
- Due Wednesday/Friday at Midnight: Labs
 - Graded “completion”-ish

How this class will run: What you'll do

- On Tuesday/Thursday: **Live Lecture**
 - Quizzes graded for completion (Polls)
 - Active participation (answer if called on)
 - In-class work (group optional; determined before class)
- Due Wednesday/Friday at Midnight: Labs
 - Graded “completion”-ish
- Due Periodically on Friday: Homework
 - About 3-4 assignments

How this class will run: What you'll do

- On Tuesday/Thursday: **Live Lecture**
 - Quizzes graded for completion (Polls)
 - Active participation (answer if called on)
 - In-class work (group optional; determined before class)
- Due Wednesday/Friday at Midnight: Labs
 - Graded “completion”-ish
- Due Periodically on Friday: Homework
 - About 3-4 assignments
- July 20th/August 12th: Final Project Dates
 - Description Presentation, Final Presentation
 - Presentations, GitHub
 - Group size: 2

How this class will run: Assessment

- Professionalism: 20%
 - Active Participation in Class/Quizzes
 - Respect your classmates, MSDS faculty, and any guests
- Assignments: 30%
 - Late Policy: 10% reduction in score per day late
 - May include Quizzes
- Labs: 20%
 - Late Policy: half credit if received less than one week late, zero otherwise
- Final Project: 30%
 - Description Presentation: 10%

How this class will run: Assessment

- Professionalism: 20%
 - Active Participation in Class/Quizzes
 - Respect your classmates, MSDS faculty, and any guests
- Assignments: 30%
 - Late Policy: 10% reduction in score per day late
 - May include Quizzes
- Labs: 20%
 - Late Policy: half credit if received less than one week late, zero otherwise
- Final Project: 30%
 - Description Presentation: 10%
- Important: Valid excuses communicated timely will be honored

How this class will run: Resources

- Office Hours
 - Tuesdays/Thursdays 3pm - 4pm
 - Online office hours
- Slack Channel
- Other MSDS faculty
- At some point (TBD)
 - In-person working sessions (maybe Wednesday or Friday)

How this class will run: Final Project

- Goal: Evidence that you are capable of utilizing deep learning to solve a task with real world data
 - Complement your existing work!
 - Get creative!
- Ideas/Starting Points
 - Kaggle Competitions (Active and Past)
 - <https://www.kaggle.com/getting-started/16221>
 - Inspiration here for [images](#) and [text](#)
 - Implement a Research Paper (very [valuable](#)!)
 - Generate Images/Text (specific categories, think Dr. Interian's recipes)
 - [More inspiration](#)

How this class will run

- Any Questions?

What is Deep Learning?

- **Deep Learning (DL)** is a subset of Machine Learning where algorithms perform tasks by extracting *high-level features* from datasets that are usually very large and unstructured.



Why do you think this is a cat?

What is Deep Learning?

- **Deep Learning (DL)** is a subset of Machine Learning where algorithms perform tasks by extracting *high-level features* from datasets that are usually very large and unstructured.



High Level Features

- Two Ears
- Two Eyes
- Whiskers
- Looks fluffy

What is Deep Learning?

- **Deep Learning (DL)** is a subset of Machine Learning where algorithms perform tasks by extracting *high-level features* from datasets that are usually very large and unstructured.



High Level Features

- Two Ears
- Two Eyes
- Whiskers
- Looks fluffy

Collectively: a representation of the image

What is Deep Learning?

- **Deep Learning (DL)** is a subset of Machine Learning where algorithms perform tasks by extracting *high-level features* from datasets that are usually very large and unstructured.



$$\begin{array}{c} \rightarrow \left[\begin{array}{ccccc} 0 & 0 & 0 & \dots & 0 \\ .5 & .75 & 1 & \dots & .25 \\ \vdots & \vdots & \vdots & & \vdots \\ .333 & 0 & 1 & \dots & 0 \end{array} \right] \begin{array}{l} \text{Two Ears} = ??? \\ \text{Two Eyes} = ??? \\ \text{Whiskers} = ??? \\ \text{Looks fluffy} = ??? \end{array} \end{array}$$

What is Deep Learning?

- **Deep Learning (DL)** is a subset of Machine Learning where algorithms perform tasks by extracting *high-level features* from datasets that are usually very large and unstructured.



Low-level geometric Features

- Edge detection
- Noisiness
- Blob detection

$$\begin{bmatrix} 0 & 0 & 0 & \dots & 0 \\ .5 & .75 & 1 & \dots & .25 \\ \vdots & \vdots & \vdots & & \vdots \\ .333 & 0 & 1 & \dots & 0 \end{bmatrix}$$

What is Deep Learning?

- **Deep Learning (DL)** is a subset of Machine Learning where algorithms perform tasks by extracting *high-level features* from datasets that are usually very large and unstructured.



DL high-level features

- learned from data
- Constructed from learned low-level features
- Usually NOT interpretable

$$\begin{bmatrix} 0 & 0 & 0 & \dots & 0 \\ .5 & .75 & 1 & \dots & .25 \\ \vdots & \vdots & \vdots & & \vdots \\ .333 & 0 & 1 & \dots & 0 \end{bmatrix}$$

What is Deep Learning?

- **Deep Learning (DL)** is a subset of Machine Learning where algorithms perform tasks by extracting *high-level features* from datasets that are usually very large and unstructured.

DL high-level features

- Usually NOT interpretable
- Only as good as your data...




Is this a snow leopard or regular leopard?

What is Deep Learning?

- **Deep Learning (DL)** is a subset of Machine Learning where algorithms perform tasks by extracting *high-level features* from datasets that are usually very large and unstructured.

DL high-level features

- Usually NOT interpretable
- Only as good as your data...




Is this a snow leopard or regular leopard?

What is Deep Learning?

- **Deep Learning (DL)** is a subset of Machine Learning where algorithms perform tasks by extracting *high-level features* from datasets that are usually very large and unstructured.

DL high-level features

- Usually NOT interpretable
- Only as good as your data...






Is this a snow leopard or regular leopard?

What is Deep Learning?

- One-Hot Embeddings -> Word Embeddings
- Unstructured data -> Represented by meaningful features
- Simple linear function

$$\begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} \quad \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix} \quad \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}$$

brain head train




What is Deep Learning?

- One-Hot Embeddings -> Word Embeddings
- Unstructured data -> Represented by meaningful features
- Simple linear function


$$\begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} \quad \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix} \quad \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}$$

brain head train



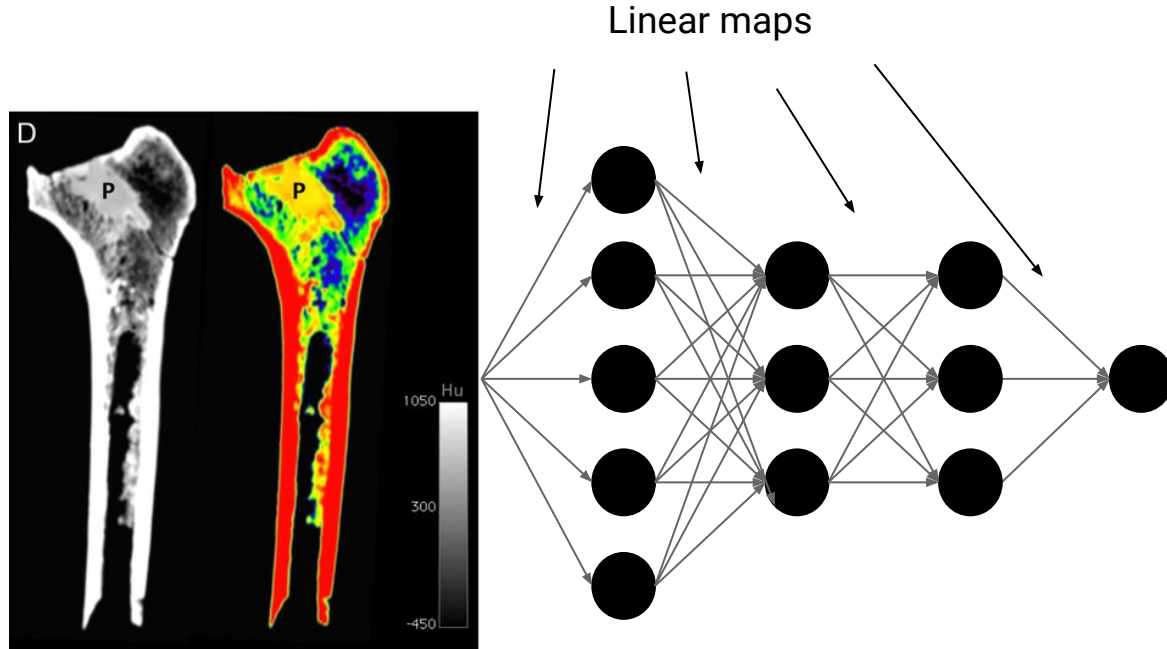
What is Deep Learning?

- Array -> Representation by “meaningful” features
- Simple linear functions stacked on top of each other




What is Deep Learning?

- Array -> Representation by “meaningful” features
- Simple linear functions stacked on top of each other





What is Deep Learning?

- Array \rightarrow Representation by “meaningful” features
- Simple linear functions stacked on top of each other




$256^2 = 65536$ dimensions




Neural Network

- Word embedding: simple linear mapping
- Neural Network: stack linear functions one after the other (layers)
- Any function can be approximated this way (rigorous!)



Neural Network

- Word embedding: simple linear mapping
- Neural Network: stack linear functions one after the other (layers)
- Any function can be approximated this way (rigorous!)
- Idea: create low level features in early layers to create high level features in later layers (none of these are necessarily interpretable!)



Neural Network

- Training a Neural Network: just a supervised learning problem!

Neural Network

- Training a Neural Network: just a supervised learning problem!
- Labelled data: $(x^{(1)}, y^{(1)}), \dots, (x^{(N)}, y^{(N)})$
- Model and parameters: $F(x; \theta)$
- Loss Function: $\mathcal{L}(F(x; \theta), y)$

Neural Network

- Training a Neural Network: just a supervised learning problem!
- Labelled data: $(x^{(1)}, y^{(1)}), \dots, (x^{(N)}, y^{(N)})$
- Model and parameters: $F(x; \theta)$
- Loss Function: $\mathcal{L}(F(x; \theta), y)$

Find parameters θ that minimize


$$\frac{1}{N} \sum_{i=1}^N \mathcal{L} \left(F(x^{(i)}; \theta), y^{(i)} \right)$$

Neural Network

- What's so special about using a Neural Network for $F(x; \theta)$?

Neural Network

- What's so special about using a Neural Network for $F(x; \theta)$?




Input data: 5D vector
(Input layer)

$$\vec{x} = \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \\ x_5 \end{bmatrix}$$

Neural Network


- First step: linear transformation



$$W^{[1]} = ?$$

Neural Network

- First step: linear transformation



Input data: 5D vector
(Input layer)


$$W^{[1]} = \begin{bmatrix} w_{11}^{[1]} & w_{12}^{[1]} & w_{13}^{[1]} & w_{14}^{[1]} & w_{15}^{[1]} \\ w_{21}^{[1]} & w_{22}^{[1]} & w_{23}^{[1]} & w_{24}^{[1]} & w_{25}^{[1]} \\ w_{31}^{[1]} & w_{32}^{[1]} & w_{33}^{[1]} & w_{34}^{[1]} & w_{35}^{[1]} \end{bmatrix}$$

$$b^{[1]} = \begin{bmatrix} b_1^{[1]} \\ b_2^{[1]} \\ b_3^{[1]} \end{bmatrix}$$

$$W^{[1]}x + b^{[1]}$$


Neural Network

- Second step: nonlinearity (activation function)



Neural Network

- Second step: nonlinearity (activation function)




Commonly used activation functions

- ReLU, Sigmoid, Leaky ReLU, Tanh

$$a^{[1]} = h(W^{[1]}x + b^{[1]})$$

Neural Network

- Let's do it again!




$$W^{[2]} = \begin{bmatrix} w_{11}^{[2]} & w_{12}^{[2]} & w_{13}^{[2]} \\ w_{21}^{[2]} & w_{22}^{[2]} & w_{23}^{[2]} \\ w_{31}^{[2]} & w_{32}^{[2]} & w_{33}^{[2]} \end{bmatrix} \quad b^{[2]} = \begin{bmatrix} b_1^{[2]} \\ b_2^{[2]} \\ b_3^{[2]} \end{bmatrix}$$

$$a^{[2]} = h(W^{[2]} a^{[1]} + b^{[2]})$$

$$a^{[1]} = h(W^{[1]} x + b^{[1]})$$

Neural Network

- Final Output




$$W^{[3]} = \begin{bmatrix} w_{11}^{[3]} & w_{12}^{[3]} & w_{13}^{[3]} \end{bmatrix} \quad b^{[3]} = \begin{bmatrix} b_1^{[3]} \end{bmatrix}$$

$$F(x; \theta) = h(W^{[3]} a^{[2]} + b^{[3]})$$

$$a^{[1]} = h(W^{[1]} x + b^{[1]})$$
$$a^{[2]} = h(W^{[2]} a^{[1]} + b^{[2]})$$

Neural Network

- 3-layer “Feed-Forward” Neural Network



$$a^{[1]} = h(W^{[1]}x + b^{[1]})$$

$$a^{[2]} = h(W^{[2]}a^{[1]} + b^{[2]})$$

$$F(x; \theta) = h(W^{[3]}a^{[2]} + b^{[3]})$$

“Deeper” = more layers

Neural Networks

- All architectures are, at their core, linearity + nonlinearity successively
 - Easy to compute gradient this way (chain rule)
- In theory all you need is a FF NN
- In practice, intuition, experience, and understanding the problem are needed to make NNs work effectively
- We will not learn every architecture under the sun (boring)

Neural Networks

- All architectures are, at their core, linearity + nonlinearity successively
 - Easy to compute gradient this way (chain rule)
- In theory all you need is a FF NN
- In practice, intuition, experience, and understanding the problem are needed to make NNs work effectively
- We will not learn every architecture under the sun (boring)

Let's do this in PyTorch!