

Deep Learning

MSDS 631

Images and
Convolutional Neural Networks

Michael Ruddy

Questions?

- From last lecture?
- From the lab assignment?

Overview



- Why imaging? Why not FF NNs?
- What/Why is a Convolution?
- CNN-specific hyperparameters
- Basic CNN history/set-up

Why Imaging?

- Humans are really good at looking at things
 - The human eye/brain is an incredibly complicated piece of machinery
- Efforts to recreate vision based on human models of vision have largely been unsuccessful

Why Imaging?

- Humans are really good at looking at things
 - The human eye/brain is an incredibly complicated piece of machinery
- Efforts to recreate vision based on human models of vision have largely been unsuccessful
- Imaging is important!



Why Imaging?

- Humans are really good at looking at things
 - The human eye/brain is an incredibly complicated piece of machinery
- Efforts to recreate vision based on human models of vision have largely been unsuccessful
- Imaging is important!
 - Classification (facial/object recognition, avoid poisonous plants, etc.)
 - Medical Imaging (detecting disease, predicting outcomes of radiation, segmentation of medical images)
 - Autonomous Driving (driver assistance, fully autonomous vehicles)
 - Deepfakes and deepfake detection

Why Imaging?

- Humans are really good at looking at things
 - The human eye/brain is an incredibly complicated piece of machinery
- Efforts to recreate vision based on human models of vision have largely been unsuccessful
- Imaging is important!
 - Classification (facial/object recognition, avoid poisonous plants, etc.)
 - Medical Imaging (detecting disease, predicting outcomes of radiation, segmentation of medical images)
 - Autonomous Driving (driver assistance, fully autonomous vehicles)
 - Deepfakes and deepfake detection
- A lot of these are time-consuming things that human can do really well

Why are images special?

- Images are deceptively hard



Why are images special?

- Images are deceptively hard

This is a cat



Why are images special?

- Images are deceptively hard

This is ??????

$$\begin{bmatrix} 0 & 0 & 0 & \dots & 0 \\ .5 & .75 & 1 & \dots & .25 \\ \vdots & \vdots & \vdots & & \vdots \\ .333 & 0 & 1 & \dots & 0 \end{bmatrix}$$

Why are images special?

- Images are deceptively hard
- Images are big



32x32 image
1024 features



512x512 image
262,144 features

Why are images special?

- Images are deceptively hard
- Images are big



32x32 image
1024 features



512x512 image
262,144 features

Fully Connected Layer

- $1024 \rightarrow 1024$
- $1024^2 = 1,048,576$
parameters
- $262,144 \rightarrow 262,144$
- $68,719,476,736$
parameters

Why are images special?

- Images are deceptively hard
- Images are big
- Geometry matters!
 - Pixels near each other interact in different ways to create features than pixels far away

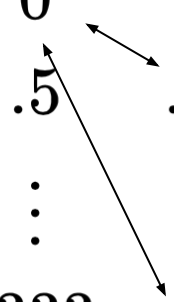
Different
relationships

$$\begin{bmatrix} 0 & 0 & 0 & \dots & 0 \\ .5 & .75 & 1 & \dots & .25 \\ \vdots & \vdots & \vdots & & \vdots \\ .333 & 0 & 1 & \dots & 0 \end{bmatrix}$$

Why are images special?

- Images are deceptively hard
- Images are big
- Geometry matters!
 - Pixels near each other interact in different ways to create features than pixels far away
 - This is free data that we lose if we simply consider an image as a data vector

Different
relationships

$$\begin{bmatrix} 0 & 0 & 0 & \dots & 0 \\ .5 & .75 & 1 & \dots & .25 \\ \vdots & \vdots & \vdots & & \vdots \\ .333 & 0 & 1 & \dots & 0 \end{bmatrix}$$


The Convolution

- Fancy **linear** operation useful for spatial data

The Convolution

- Fancy **linear** operation useful for spatial data

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

The Convolution

- Fancy **linear** operation useful for spatial data

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

*

The Convolution

- Fancy **linear** operation useful for spatial data

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

*

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

=

$$\begin{bmatrix} & & & \\ & & & \\ & & & \\ & & & \end{bmatrix}$$

The Convolution

- Fancy **linear** operation useful for spatial data
- Element-wise product

$$(1 \times 1) + (.5 \times 0) + (0 \times 0) + (.25 \times 2) \\ = 1.5$$

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

*

$$= \begin{bmatrix} 1.5 & \dots & \dots \\ \vdots & \ddots & \vdots \end{bmatrix}$$

The Convolution

- Fancy **linear** operation useful for spatial data
- Element-wise product

$$(.5 \times 1) + (1 \times 0) + (.25 \times 0) + (.5 \times 2) = 1.5$$

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

*

$$= \begin{bmatrix} 1.5 & \boxed{1.5} \end{bmatrix}$$

The Convolution

- Fancy **linear** operation useful for spatial data
- Element-wise product

$$(1 \times 1) + (0 \times 0) + (.5 \times 0) + (1 \times 2) = 3$$

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

*

$$= \begin{bmatrix} 1.5 & 1.5 & \rightarrow 3 \end{bmatrix}$$

The Convolution

- Fancy **linear** operation useful for spatial data
- Element-wise product

$$(0 \times 1) + (.25 \times 0) + (1 \times 0) + (.25 \times 2) = .5$$

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

*



$$\begin{bmatrix} 1.5 & 1.5 & 3 \\ 0.5 & & \end{bmatrix}$$

The Convolution

- Fancy **linear** operation useful for spatial data
- Element-wise product

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

*

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

=

$$\begin{bmatrix} 1.5 & 1.5 & 3 \\ 0.5 & ? & ? \\ ? & ? & ? \end{bmatrix}$$

The Convolution

- Fancy **linear** operation useful for spatial data
- Element-wise product

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

*

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

=

$$\begin{bmatrix} 1.5 & 1.5 & 3 \\ 0.5 & .25 & 2.5 \\ 1 & 2.25 & 2 \end{bmatrix}$$

The Convolution

- Fancy **linear** operation useful for spatial data
- Element-wise product

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} w_{11}^{(1)} & w_{12}^{(1)} \\ w_{21}^{(1)} & w_{22}^{(1)} \end{bmatrix} = \text{[Output Image]}$$

Unknown Parameters

“2x2 Filter”

Why Convolution?

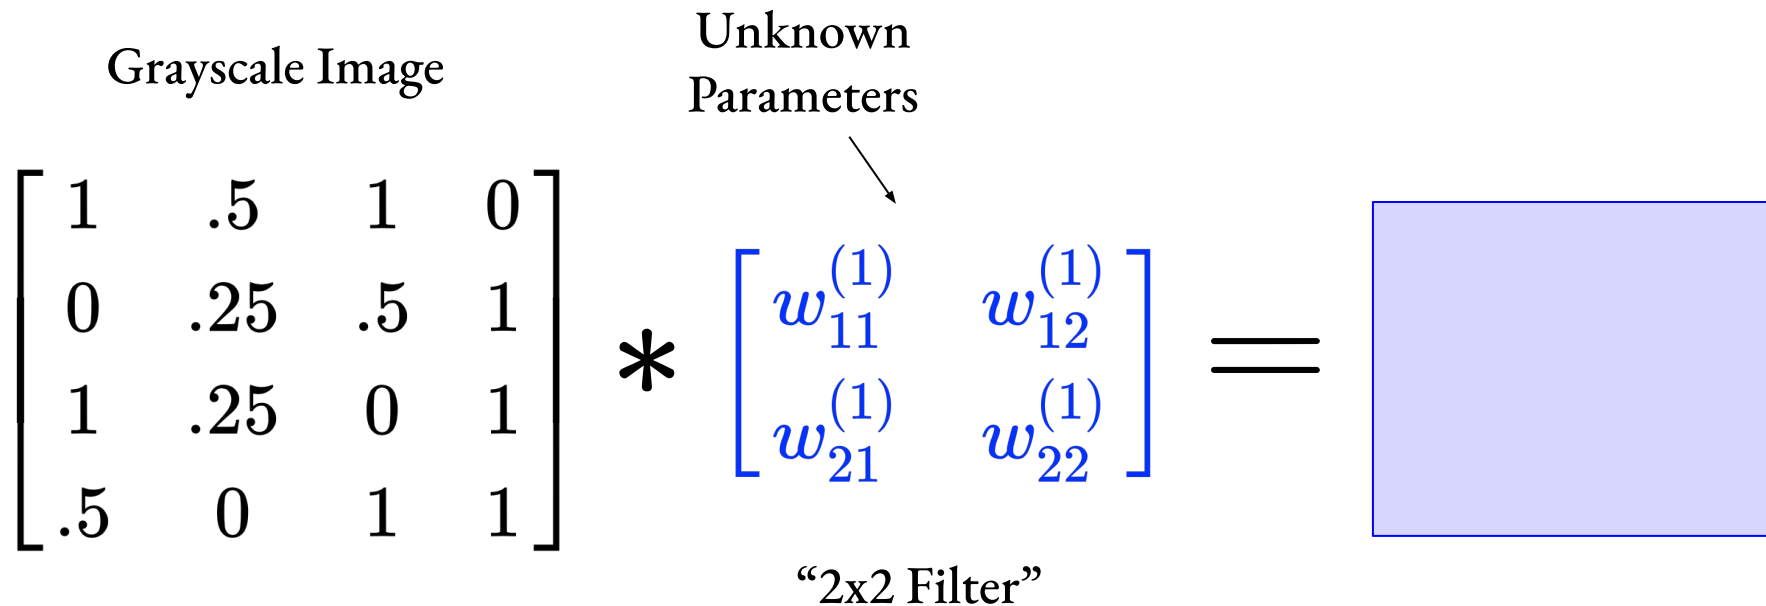
- Only four parameters!
 - If input is dimension 16 and output is dimension 9, how many for FC?

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} w_{11}^{(1)} & w_{12}^{(1)} \\ w_{21}^{(1)} & w_{22}^{(1)} \end{bmatrix} = \text{Output Image}$$

Unknown Parameters

“2x2 Filter”



Why Convolution?


- Only four parameters!
- Translational Equivariance
 - If I shift my image, I shift the output!

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} w_{11}^{(1)} & w_{12}^{(1)} \\ w_{21}^{(1)} & w_{22}^{(1)} \end{bmatrix} = \text{Image}$$

Unknown Parameters

“2x2 Filter”



Why Convolution?

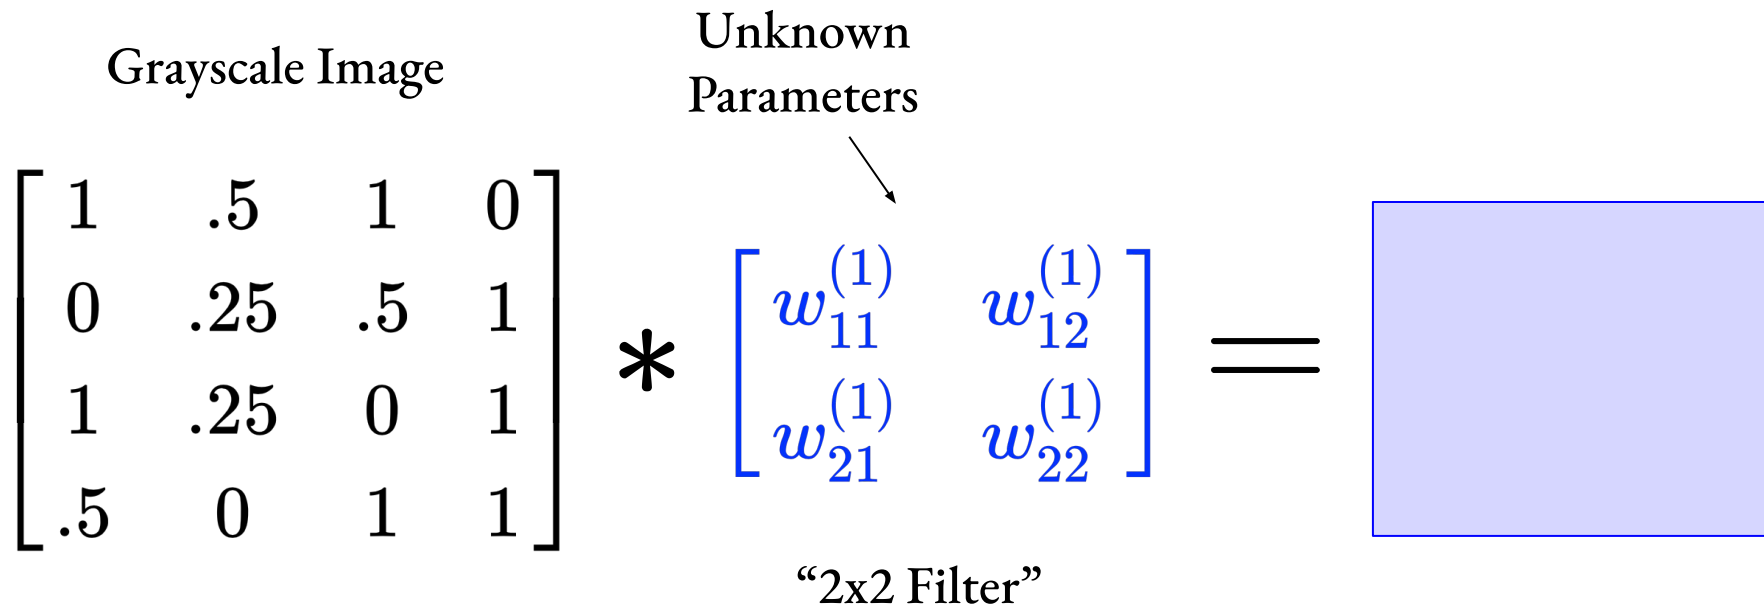
- Only four parameters!
- Translational Equivariance
- Weight Sharing (detect same feature translated to different parts of the image)

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} w_{11}^{(1)} & w_{12}^{(1)} \\ w_{21}^{(1)} & w_{22}^{(1)} \end{bmatrix} = \text{[Output Box]}$$

Unknown Parameters

“2x2 Filter”



Why Convolution?

Intuition: Edge
Detection

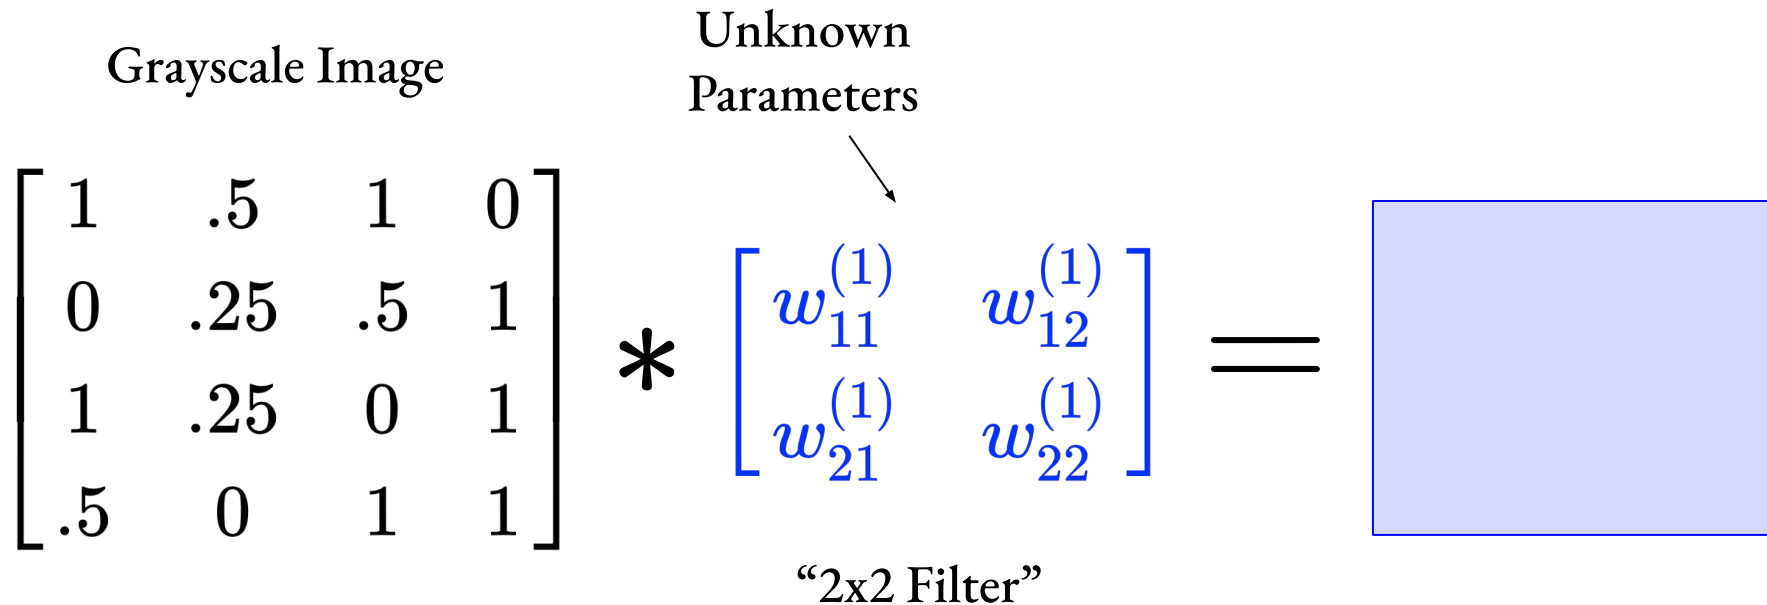
- Only four parameters!
- Translational Equivariance
- Weight Sharing (detect same feature translated to different parts of the image)

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} w_{11}^{(1)} & w_{12}^{(1)} \\ w_{21}^{(1)} & w_{22}^{(1)} \end{bmatrix} = \text{[Output Box]}$$

Unknown Parameters

“2x2 Filter”



The Convolution

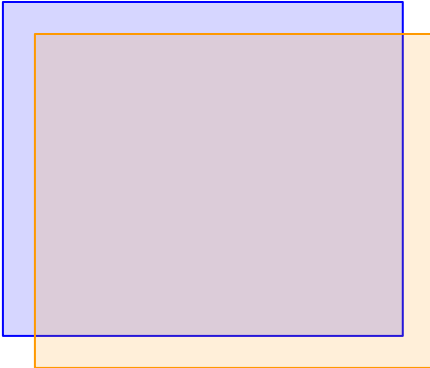
- In a Conv. layer we apply many filter to get many features

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} w_{11}^{(2)} & w_{12}^{(2)} \\ w_{21}^{(2)} & w_{22}^{(2)} \end{bmatrix} =$$

More Parameters

“2x2 Filter”



The Convolution

- In a Conv. layer we apply many filter to get many features
- Applying N filters to an image results in an output with N “channels”

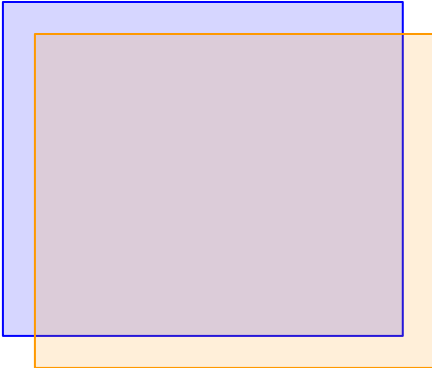
Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} w_{11}^{(2)} & w_{12}^{(2)} \\ w_{21}^{(2)} & w_{22}^{(2)} \end{bmatrix} =$$

More Parameters


“2x2 Filter”

“2 Channels”





Layer 3



Layer 2



Layer 1

The Convolution

- In a Conv. layer we apply many filter to get many features
- Applying N filters to an image results in an output with N “channels”

RGB Image

$$\begin{bmatrix} \begin{bmatrix} 1 \\ 0 \\ 1 \\ .5 \\ 0 \end{bmatrix} \begin{bmatrix} .2 \\ .2 \\ 1 \\ 1 \\ 0 \end{bmatrix} \begin{bmatrix} .5 \\ .25 \\ .25 \\ 10 \\ 0 \end{bmatrix} \begin{bmatrix} .15 \\ .25 \\ 1 \\ .5 \\ 0 \end{bmatrix} \begin{bmatrix} 1 \\ .5 \\ 10 \\ .25 \\ .75 \end{bmatrix} \begin{bmatrix} 0 \\ 1 \\ .5 \\ 1 \\ 0 \end{bmatrix} \begin{bmatrix} 0 \\ 0 \\ 1 \\ .2 \\ 0 \end{bmatrix} \begin{bmatrix} 0 \\ 1 \\ 1 \\ 1 \\ 0 \end{bmatrix} \end{bmatrix}$$

The Convolution


- In a Conv. layer we apply many filter to get many features
- Applying N filters to an image results in an output with N “channels”

RGB Image $4 \times 4 \times 3$

$$\begin{bmatrix} \begin{bmatrix} 1 \\ 0 \\ 1 \\ .5 \end{bmatrix} \begin{bmatrix} .2 \\ .2 \\ 1 \\ 0 \end{bmatrix} \begin{bmatrix} .5 \\ .25 \\ .25 \\ 1 \end{bmatrix} \begin{bmatrix} .15 \\ .25 \\ 1 \\ 0 \end{bmatrix} \begin{bmatrix} 1 \\ .5 \\ 1 \\ .25 \end{bmatrix} \begin{bmatrix} 0 \\ 1 \\ .5 \\ 0 \end{bmatrix} \begin{bmatrix} 0 \\ 0 \\ 1 \\ .2 \end{bmatrix} \end{bmatrix}$$


The Convolution

- In a Conv. layer we apply many filter to get many features
- Applying N filters to an image results in an output with N “channels”




The Convolution

- In a Conv. layer we apply many filter to get many features
- Applying N filters to an image results in an output with N “channels”




The Convolution

- In a Conv. layer we apply many filter to get many features
- Applying N filters to an image results in an output with N “channels”
- Filter channels must match input channels!!!




The Convolution

- In a Conv. layer we apply many filter to get many features
- Applying N filters to an image results in an output with N “channels”
- Filter channels must match input channels!!!




The Convolution

- In a Conv. layer we apply many filter to get many features
- Applying N filters to an image results in an output with N “channels”
- Filter channels must match input channels!!!




The Convolution

- In a Conv. layer we apply many filter to get many features
- Applying N filters to an image results in an output with N “channels”
- Filter channels must match input channels!!!




The Convolution

- In a Conv. layer we apply many filter to get many features
- Applying N filters to an image results in an output with N “channels”
- Filter channels must match input channels!!!



The Convolution



Convolution Hyperparameters

- Number of Filters

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
 - “How far it jumps when sliding”

Stride 1

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
 - “How far it jumps when sliding”

Stride 1

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
 - “How far it jumps when sliding”

Stride 1

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
 - “How far it jumps when sliding”

Stride 1

The diagram illustrates a 1D convolution operation. On the left is a 4x4 input matrix with values: $\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$. A 2x2 region of this matrix is highlighted with a red dashed border, containing the values $\begin{bmatrix} 0 & .25 \\ 1 & .25 \end{bmatrix}$. To the right of this region is a multiplication symbol $*$. Further right is a 2x2 kernel matrix highlighted with a blue dashed border, containing the values $\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$. Blue dashed lines connect the corners of the red dashed box to the corners of the blue dashed box, indicating the alignment for the convolution operation. The text "Stride 1" is positioned above the kernel matrix.

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
 - “How far it jumps when sliding”


Stride 2

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
 - “How far it jumps when sliding”

Stride 2



Convolution Hyperparameters

- Number of Filters
- Stride of the filter
 - “How far it jumps when sliding”

Stride 2

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
 - What is the dimension of the output for Stride 1 vs. Stride 2?

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
- Size of filter

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 5 & 1 \\ 0 & 1 & 2 \\ 1 & 1 & 0 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
- Size of filter
 - What is output dimension here if stride = 1?

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 5 & 1 \\ 0 & 1 & 2 \\ 1 & 1 & 0 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
- Size of filter
- Problem: size of output keep shrinking!
 - Only a few convolutional layers before the resulting 2D dimensions are very small

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
- Size of filter
- Problem: size of output keep shrinking!
 - Only a few convolutional layers before the resulting 2D dimensions are very small
- Solution: Zero padding

$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & .5 & 1 & 0 & 0 \\ 0 & 0 & .25 & .5 & 1 & 0 \\ 0 & 1 & .25 & 0 & 1 & 0 \\ 0 & .5 & 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

Convolution Hyperparameters

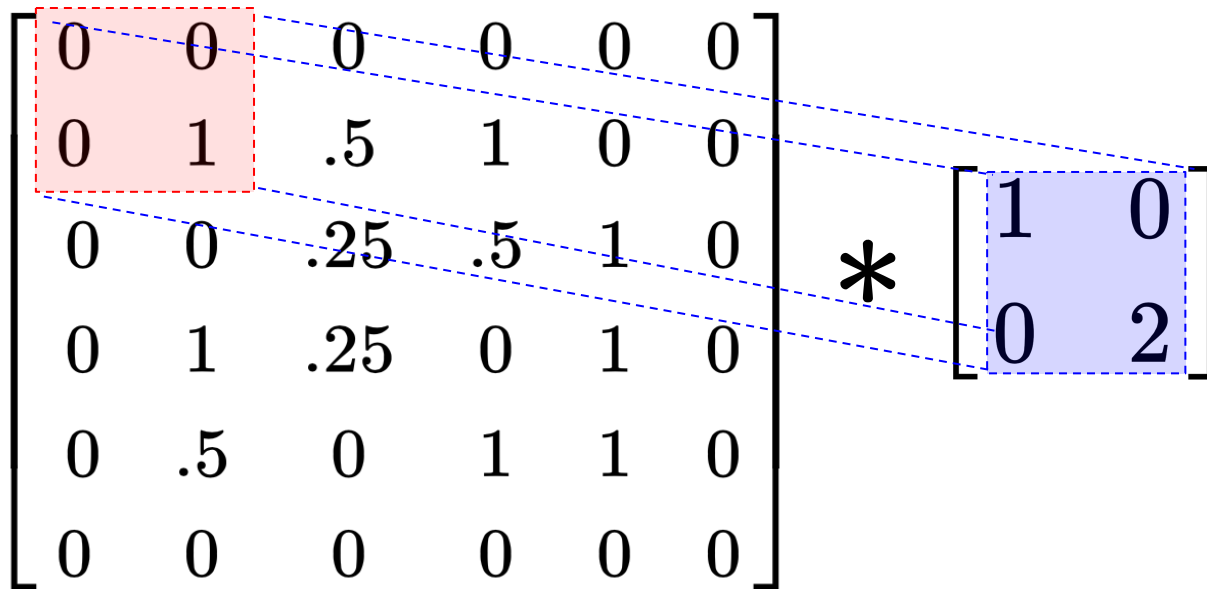
- Number of Filters
- Stride of the filter
- Size of filter
- Padding

$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & .5 & 1 & 0 & 0 \\ 0 & 0 & .25 & .5 & 1 & 0 \\ 0 & 1 & .25 & 0 & 1 & 0 \\ 0 & .5 & 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

Convolution Hyperparameters

- Number of Filters
- Stride of the filter
- Size of filter
- Padding

Padding
by one



The diagram illustrates a 1D convolution operation. On the left is a 6x6 input matrix with a red dashed box highlighting the top-left 2x2 region (values 0, 0, 0, 1). This region is connected by blue dashed lines to a 2x2 kernel matrix on the right (values 1, 0, 0, 2), which is highlighted with a blue dashed box. A black asterisk (*) is placed between the two matrices. An arrow points from the text 'Padding by one' to the top-right corner of the input matrix, indicating the padding applied to the original 4x4 data.

$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & .5 & 1 & 0 & 0 \\ 0 & 0 & .25 & .5 & 1 & 0 \\ 0 & 1 & .25 & 0 & 1 & 0 \\ 0 & .5 & 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

Convolution Hyperparameters

- Common choices for a Conv-Layer:
 - Stride = 1
 - Odd Filter Size (3x3, 5x5, etc.)
 - “Same” padding

Diagram illustrating a 1D convolution operation with a 6x6 input, a 3x3 kernel, and a 4x4 output. The input is a 6x6 grid with a 3x3 region highlighted in red. The kernel is a 3x3 grid highlighted in blue. The output is a 4x4 grid. Blue dashed lines connect the corners of the red region to the corners of the blue region, and from the corners of the blue region to the corners of the output grid. A multiplication symbol (*) is placed between the input and kernel grids.

0	0	0	0	0	0
0	1	.5	1	0	0
0	0	.25	.5	1	0
0	1	.25	0	1	0
0	.5	0	1	1	0
0	0	0	0	0	0

*

1	0	1
0	1	1
1	.5	2


Convolution Hyperparameters

- Common choices for a Conv-Layer:
 - Stride = 1
 - Odd Filter Size (3x3, 5x5, etc.)
 - “Same” padding

$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & .5 & 1 & 0 & 0 \\ 0 & 0 & .25 & .5 & 1 & 0 \\ 0 & 1 & .25 & 0 & 1 & 0 \\ 0 & .5 & 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} * \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 1 \\ 1 & .5 & 2 \end{bmatrix}$$


Convolutional Neural Network

- LeNet-5 by Yann LeCun




Convolutional Neural Network

- LeNet-5 by Yann LeCun




Convolutional Neural Network

- LeNet-5 by Yann LeCun




Convolutional Neural Network

- LeNet-5 by Yann LeCun



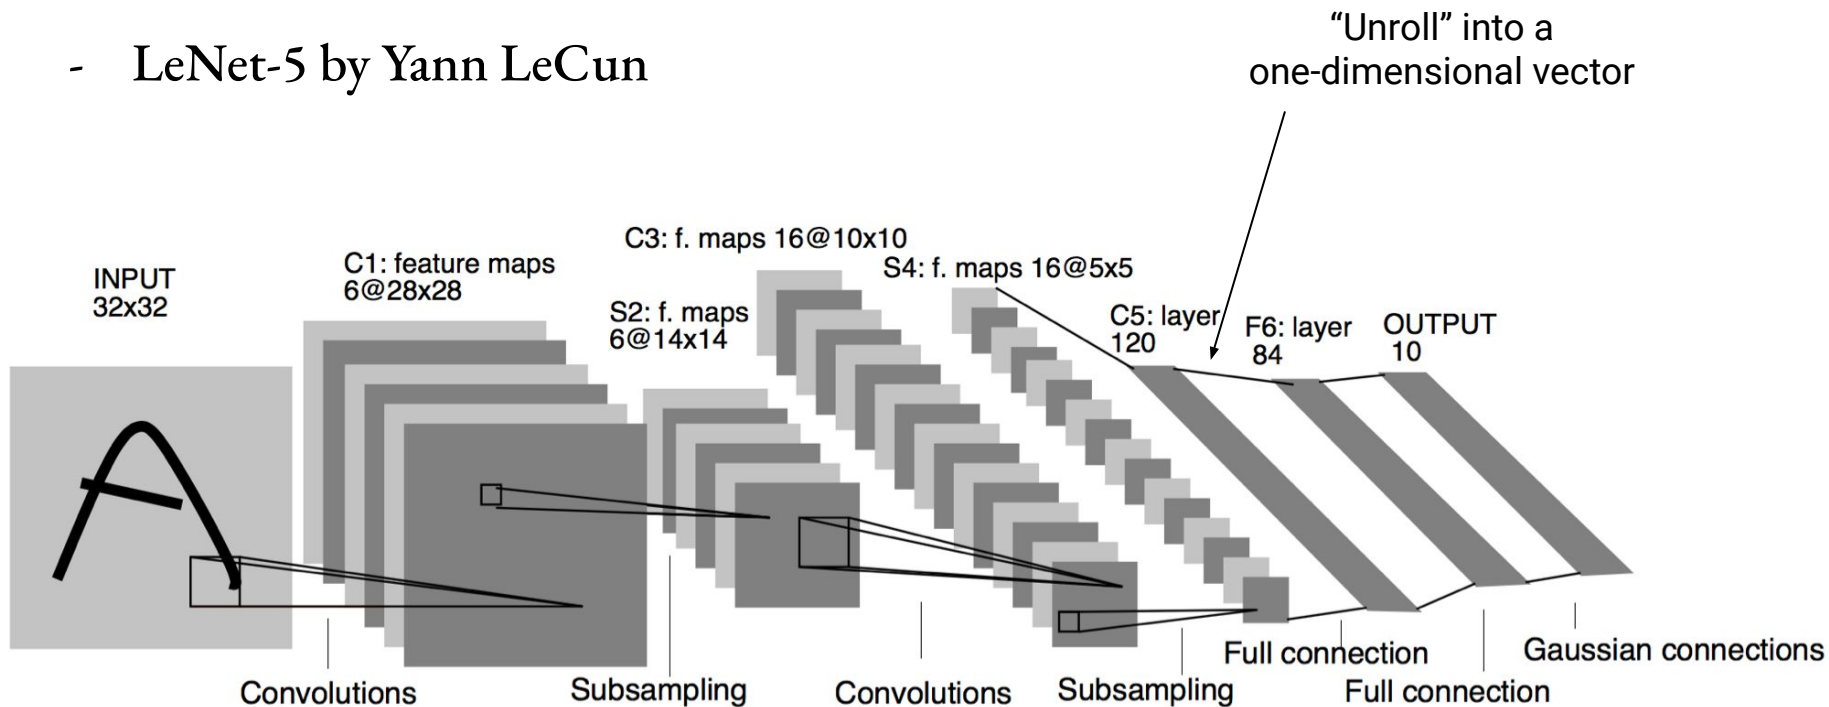
Convolutional Neural Network

- LeNet-5 by Yann LeCun



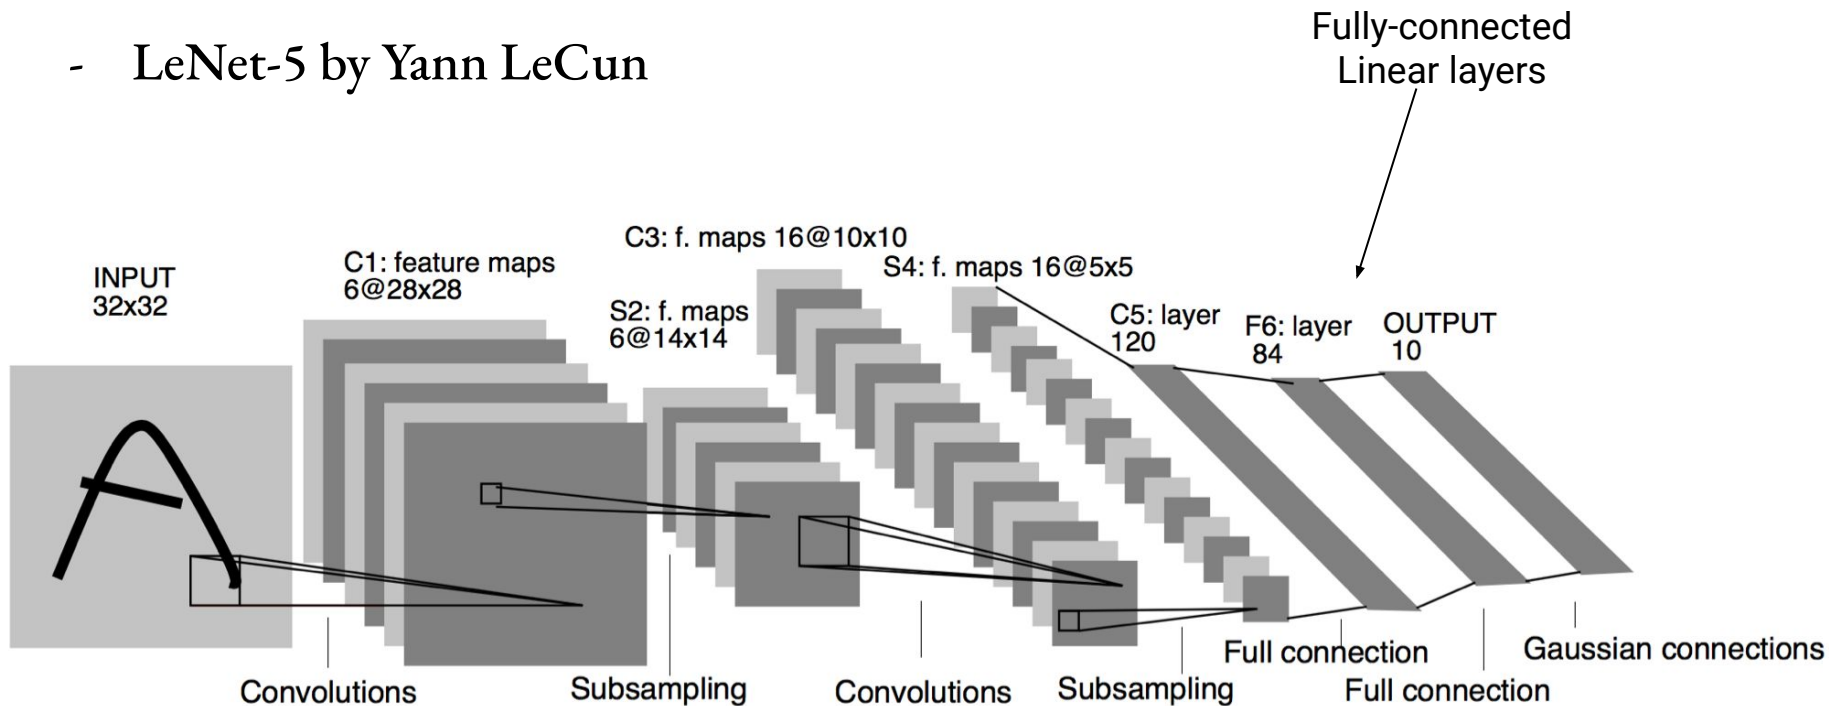
Convolutional Neural Network

- LeNet-5 by Yann LeCun



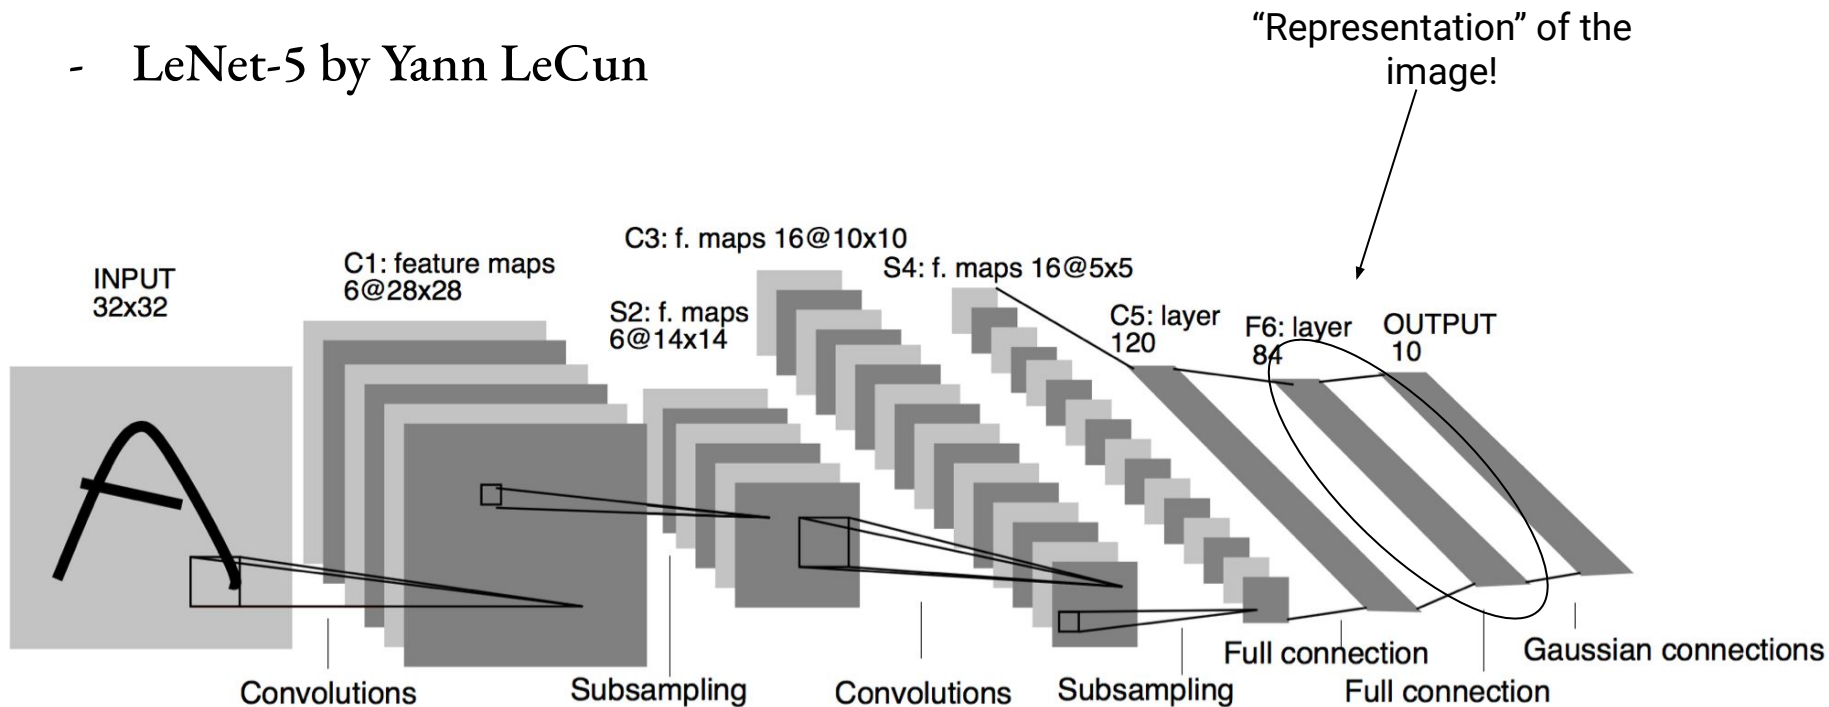
Convolutional Neural Network

- LeNet-5 by Yann LeCun



Convolutional Neural Network

- LeNet-5 by Yann LeCun



Convolutional Neural Network

- AlexNet wins ImageNet Competition in 2012




Figure 2: An illustration of the architecture of our CNN, explicitly showing the delineation of responsibilities between the two GPUs. One GPU runs the layer-parts at the top of the figure while the other runs the layer-parts at the bottom. The GPUs communicate only at certain layers. The network's input is 150,528-dimensional, and the number of neurons in the network's remaining layers is given by 253,440–186,624–64,896–64,896–43,264–4096–4096–1000.

Convolutional Neural Network

- AlexNet wins ImageNet Competition in 2012
- By 2015 we have CNNs with >100 layers, better than human-level performance





Figure 2: An illustration of the architecture of our CNN, explicitly showing the delineation of responsibilities between the two GPUs. One GPU runs the layer-parts at the top of the figure while the other runs the layer-parts at the bottom. The GPUs communicate only at certain layers. The network's input is 150,528-dimensional, and the number of neurons in the network's remaining layers is given by 253,440–186,624–64,896–64,896–43,264–4096–4096–1000.

Downsampling

- Reduce size of output
- Minimal information loss in practice
- Intuition: reduce resolution of the image


Downsampling

- Reduce size of output
- Minimal information loss in practice
- Intuition: reduce resolution of the image
- Max Pooling




Downsampling

- Reduce size of output
 - Minimal information loss in practice
 - Intuition: reduce resolution of the image
 - Max Pooling
- 2x2 filter size
 - Stride 2




Downsampling

- Reduce size of output
 - Minimal information loss in practice
 - Intuition: reduce resolution of the image
 - Max Pooling
- 2x2 filter size
 - Stride 2




Downsampling

- Reduce size of output
 - Minimal information loss in practice
 - Intuition: reduce resolution of the image
 - Max Pooling
 - Average Pooling
- 2x2 filter size
 - Stride 2




Downsampling

- Done along spatial dimension, preserves channels



Convolutional Neural Network

- LeNet-5 by Yann LeCun



Summary

- Convolution Layers
 - Suited for Spatial Data
 - Less Parameters than FC Layers, Weight sharing
- Common Hyperparameters
 - Number of Filters, Filter Size, Stride, Padding
- Common Sequence
 - Conv -> Activation -> Conv -> Activation -> Downsampling
 - Repeat until unrolled into final FC layers